

CSCI 6730 Operating Systems

Project #1: Multi-process and IPC

Xiaodong Jiang
Department of Statistics
Feb. 6th, 2017

Introduction

In this project, we design and implement a multiprocess word counting program by converting a given single-process word counting program in C language, the key idea is use **fork()** to create child and parent processes to read and count the number of word/char/lines simultaneously, while employing **pipe()** to achieve the inter-process communication (**IPC**). In the end, we compare the running time of single-process and multiprocess versions, and observed a great advantage to use multiprocess program, which runs much faster while achieving the same results.

Main idea on fork and pipe (IPC)

In the main function, instead of implementing a sequential loop to read the data one by one, we first create a pipe, which could return two file descriptors referring to the ends of the pipe, i.e., write end and read end.

```
90 // create pipe pair
91 int fp[2];
92 if (pipe(fp) == 0) {
93     //create pipe
94     //fork children process
95     for (i = 0; i < numFiles; i++) {
96         pid = fork();
97         if (pid == -1) {
98             fprintf(stderr, "Fork failure");
99             exit(EXIT_FAILURE);
100         }
101         if (pid == 0) {
102             printf("*Child[%d] is create \n", getpid());
103             sprintf(filename, "%s/text.%02d", FILEPATH, i);
104             printf("read: %s\n", filename);
105
106             tmp = word_count(filename);
107             printf("Child[%d] result: chars count is %d, lines count is %d, words count is %d\n",
108                   getpid(), tmp.charcount, tmp.linecount, tmp.wordcount);
109             close(fp[0]);
110
111             data_processed = write(fp[1], &tmp, sizeof(tmp));
112             printf("Child[%d] is exit. \n", getpid());
113             exit(EXIT_SUCCESS);
114         }
115     }
116 }
117
118 close(fp[1]);
```

As shown in the screenshot above, we fork n pairs of processes at the read end of pipe with a loop in *line 95-116*, while first making some **error message handling** depends

on the process id (pid), then call the child function - **word_cout** - in each child process when (pid==0). An important step here, is to close the read end before collecting the results from all child process, as we did in *line 118*.

Now, we make another loop to collect the results from all child process above, see the code below,

```

118     close(fp[1]);
119
120     int j;
121     for (j = 0; j < numFiles; j++) {
122         count_t kk[1];
123         data_processed = read(fp[0], kk, sizeof(kk));
124         count.charcount += kk->charcount;
125         count.linecount += kk->linecount;
126         count.wordcount += kk->wordcount;
127     }
128     //exit(EXIT_SUCCESS);
129
130 }
131
132
133 printf("=====\n");
134 printf("Total Lines : %d \n", count.linecount);
135 printf("Total Words : %d \n", count.wordcount);
136 printf("Total Characters : %d \n", count.charcount);
137 printf("=====\n");
138
139 return (0);
140 }

```

In this loop, we create a `count_t` data structure and read the data/results from the pipe, and add three kind of count to the corresponding variables in **count**.

Results Comparison on nike Machine

We can verified that, to read 10 files, the multiprocess program is much faster (almost 10 times) while obtaining the same results.

Single-Process Word Counting	Multiprocess Word Counting
<pre> ===== Total Lines : 16177972 Total Words : 151538006 Total Characters : 665714062 ===== real 0m10.718s user 0m10.400s sys 0m0.316s </pre>	<pre> ===== Total Lines : 16177972 Total Words : 151538006 Total Characters : 665714062 ===== real 0m1.394s user 0m0.001s sys 0m0.001s </pre>